

Big Data Analytics Structuring with Regex

Muhammad Affan Alim

Regex for Structuring (Data Wrangling)

What is Feature Engineering - Introduction

- The regular expression language is relatively small and restricted, so not all possible string processing tasks can be done using regular expressions

Match Characters

- Some characters are special **metacharacters**, and don't match themselves.
- Instead, they signal that some out-of-the-ordinary thing should be matched, or they affect other portions of the RE by repeating them or changing their meaning.
- Much of this document is devoted to discussing various metacharacters and what they do.

Here's a complete list of the metacharacters

. ^ \$ * + ? { } [] \ | ()

Here's a complete list of the metacharacters

- The first metacharacters we'll look at are `[` and `]`. They're used for specifying a character class, which is a set of characters that you wish to match.
- For example, `[abc]` will match any of the characters `a`, `b`, or `c`; this is the same as `[a-c]`, which uses a range to express the same set of characters.
- If you wanted to match only lowercase letters, your RE would be `[a-z]`.

Here's a complete list of the metacharacters

- Metacharacters are not active inside classes. For example, `[akm$]` will match any of the characters 'a', 'k', 'm', or '\$'; '\$' is usually a metacharacter, but inside a character class it's stripped of its special nature.

Here's a complete list of the metacharacters

- You can match the characters not listed within the class by complementing the set.
- This is indicated by including a '^' as the first character of the class.
- For example, `[^5]` will match any character except '5'. If the caret appears elsewhere in a character class, it does not have special meaning. For example: `[5^]` will match either a '5' or a '^'.

Here's a complete list of the metacharacters

- Perhaps the most important metacharacter is the backslash, \
- As in Python string literals, the backslash can be followed by various characters to signal various special sequences. It's also used to escape all the metacharacters so you can still match them in patterns;
- for example, if you need to match a [or \, you can precede them with a backslash to remove their special meaning: \[or \\\.

Here's a complete list of the metacharacters

- Some of the special sequences beginning with \' represent predefined sets of characters that are often useful,
- Let's take an example: \w matches any alphanumeric character. If the regex pattern is expressed in bytes, this is equivalent to the class [a-zA-Z0-9_].



Here's a complete list of the metacharacters

- **\d**
Matches any decimal digit; this is equivalent to the class [0-9].
- **\D**
Matches any non-digit character; this is equivalent to the class [^0-9].
- **\s**
Matches any whitespace character; this is equivalent to the class [\t\n\r\f\v].


Here's a complete list of the metacharacters

- **\S**
Matches any non-whitespace character; this is equivalent to the class [^\t\n\r\f\v].
- **\w**
Matches any alphanumeric character; this is equivalent to the class [a-zA-Z0-9_].
- **\W**
Matches any non-alphanumeric character; this is equivalent to the class [^a-zA-Z0-9_].

Here's a complete list of the metacharacters

- These sequences can be included inside a character class. For example, `[\s,.]` is a character class that will match any whitespace character, or ',' or '!'.

- The [final metacharacter](#) in this section is `'.`. It matches anything except a newline character, and there's an alternate mode (`re.DOTALL`) where it will match even a newline. `.` is often used where you want to match "any character".


Here's a complete list of the metacharacters

- **Repeating Things**
- Kleen `*` and kleen `+`
- There are two more repeating qualifiers. The question mark character, `?`, matches either once or zero times; you can think of it as marking something as being optional. For example, `home-?brew` matches either `'homebrew'` or `'home-brew'`.


Here's a complete list of the metacharacters

- The most complicated repeated qualifier is `{m, n}`, where `m` and `n` are decimal integers.
- This qualifier means there must be at least `m` repetitions, and at most `n`.
- For example, `a/{1,3}b` will match `'a/b'`, `'a//b'`, and `'a///b'`.
- It won't match `'ab'`, which has no slashes, or `'a////b'`, which has four.

Here's a complete list of the metacharacters

- You can omit either `m` or `n`; in that case, a reasonable value is assumed for the missing value.
- Omitting `m` is interpreted as a lower limit of 0, while omitting `n` results in an upper bound of infinity.

Here's a complete list of the metacharacters

- Readers of a reductionist bent may notice that the three other qualifiers can all be expressed using this notation. `{0,}` is the same as `*`, `{1,}` is equivalent to `+`, and `{0,1}` is the same as `?`.
- It's better to use `*`, `+`, or `?` when you can, simply because they're shorter and easier to read.

Here's a complete list of the metacharacters

•

(Dot.) In the default mode, this matches any character except a newline. If the DOTALL flag has been specified, this matches any character including a newline.

`^`

(Caret.) Matches the start of the string, and in MULTILINE mode also matches immediately after each newline.

Here's a complete list of the metacharacters

\$

Matches the end of the string or just before the newline at the end of the string, and in MULTILINE mode also matches before a newline. `foo` matches both 'foo' and 'foobar', while the regular expression `foo$` matches only 'foo'. More interestingly, searching for `foo.$` in 'foo1\nfoo2\n' matches 'foo2' normally, but 'foo1' in MULTILINE mode; searching for a single `$` in 'foo\n' will find two (empty) matches: one just before the newline, and one at the end of the string.

Module Contents

- The module defines several functions, constants, and an exception. Some of the functions are simplified versions of the full featured methods for compiled regular expressions.
- Most non-trivial applications always use the compiled form.

Module Contents

- `re.compile(pattern, flags=0)`
- Compile a regular expression pattern into a regular expression object, which can be used for matching using its `match()`, `search()` and other methods, described below.

Module Contents

- The sequence

```
>> prog = re.compile(pattern)
```

```
>> result = prog.match(string)
```
- is equivalent to
- `result = re.match(pattern, string)`
- but using `re.compile()` and saving the resulting regular expression object for reuse is more efficient when the expression will be used several times in a single program.

Module Contents

- **Performing Matches**

Once you have an object representing a compiled regular expression, what do you do with it? Pattern objects have several methods and attributes. Only the most significant ones will be covered here; consult the re docs for a complete listing.

Module Contents

Method/Attribute	Purpose
<code>match()</code>	Determine if the RE matches at the beginning of the string.
<code>search()</code>	Scan through a string, looking for any location where this RE matches.
<code>findall()</code>	Find all substrings where the RE matches, and returns them as a list.
<code>finditer()</code>	Find all substrings where the RE matches, and returns them as an iterator .

Regex Function

- The `re` module offers a set of functions that allows us to search a string for a match:

Function	Description
<code>findall</code>	Returns a list containing all matches
<code>search</code>	Returns a <code>Match object</code> if there is a match anywhere in the string
<code>split</code>	Returns a list where the string has been split at each match
<code>sub</code>	Replaces one or many matches with a string

Module Contents

- `>>> import re`
- `>>> p = re.compile('[a-z]+')`
- `>>> p`
- `re.compile('[a-z]+')`

Module Contents

- Now, you can try matching various strings against the RE `[a-z]+`.
- An empty string shouldn't match at all, since `+` means 'one or more repetitions'.
- `match()` should return `None` in this case

```
>>> p.match("")
```

```
>>> print(p.match(""))
```

```
None
```

Metacharacters

- Metacharacters are characters with a special meaning:

Character	Description	Example
<code>[]</code>	A set of characters	<code>"[a-m]"</code>
<code>\</code>	Signals a special sequence (can also be used to escape special characters)	<code>"\d"</code>
<code>.</code>	Any character (except newline character)	<code>"he..o"</code>
<code>^</code>	Starts with	<code>"^hello"</code>
<code>\$</code>	Ends with	<code>"planet\$"</code>
<code>*</code>	Zero or more occurrences	<code>"he.*o"</code>
<code>+</code>	One or more occurrences	<code>"he.+o"</code>
<code>?</code>	Zero or one occurrences	<code>"he.?o"</code>
<code>{}</code>	Exactly the specified number of occurrences	<code>"he{2}o"</code>
<code> </code>	Either or	<code>"falls stays"</code>
<code>()</code>	Capture and group	

Metacharacters

- []

```
>> import re
```

```
>> txt = "The rain in Spain"
```

```
>> #Find all lower case characters alphabetically between "a" and  
"m":
```

```
>> x = re.findall("[a-m]", txt)
```

```
>> print(x)
```

Metacharacters

- \

```
>> import re
```

```
>> txt = "That will be 59 dollars"
```

```
>> #Find all digit characters:
```

```
>> x = re.findall("\d", txt)
```

```
>> print(x)
```

Metacharacters

- .

```
>> import re
```

```
>> txt = "hello planet"
```

```
>> #Search for a sequence that starts with "he", followed by two  
(any) characters, and an "o":
```

```
>> x = re.findall("he..o", txt)
```

```
>> print(x)
```

Output: ['hello']

Metacharacters

- ^

```
>> import re
```

```
>> txt = "hello planet"
```

```
>> #Check if the string starts with 'hello':
```

```
>> x = re.findall("^hello", txt)
```

```
>> if x:
```

```
>>     print("Yes, the string starts with 'hello'")
```

```
>> else:
```

```
>>     print("No match")
```


Metacharacters

- \$

```
>> import re
>> txt = "hello planet"
>> #Check if the string ends with 'planet':
>> x = re.findall("planet$", txt)
>> if x:
>>     print("Yes, the string ends with 'planet'")
>> else:
>>     print("No match")
```

Metacharacters

- *

```
>> import re
>> txt = "hello planet"
>> #Search for a sequence that starts with "he", followed by 0 or more (any)
>> characters, and an "o":
>> x = re.findall("he.*o", txt)
>> print(x)
```

Metacharacters

- **+**

```
>> import re
>> txt = "hello planet"
>> #Search for a sequence that starts with "he", followed by 1 or more (any)
characters, and an "o":
>> x = re.findall("he.+o", txt)
>> print(x)
```

- Output: ['hello']

Metacharacters

- **?**

```
>> import re
>> txt = "hello planet"
>> #Search for a sequence that starts with "he", followed by 0 or 1 (any)
character, and an "o":
>> x = re.findall("he.?o", txt)
>> print(x)
>> #This time we got no match, because there were not zero, not one, but two
characters between "he" and the "o"
```

Output: []

Metacharacters

- `{}`

```
>> import re
```

```
>> txt = "hello planet"
```

```
>> #Search for a sequence that starts with "he", followed exactly 2  
(any) characters, and an "o":
```

```
>> x = re.findall("he.{2}o", txt)
```

```
>> print(x)
```

Output: ['hello']

Metacharacters

- `|`

```
>> import re
```

```
>> txt = "The rain in Spain falls mainly in the plain!"
```

```
>> #Check if the string contains either "falls" or "stays":
```

```
>> x = re.findall("falls|stays", txt)
```

```
>> print(x)
```

```
>> if x:
```

```
>>     print("Yes, there is at least one match!")
```

```
>> else:
```

```
>>     print("No match")
```

Special Sequences

- A special sequence is a `\` followed by one of the characters in the list below, and has a special meaning:

Character	Description	Example
<code>\A</code>	Returns a match if the specified characters are at the beginning of the string	<code>"\AThe"</code>
<code>\b</code>	Returns a match where the specified characters are at the beginning or at the end of a word (the "r" in the beginning is making sure that the string is being treated as a "raw string")	<code>r"\bain"</code> <code>r"ain\b"</code>
<code>\B</code>	Returns a match where the specified characters are present, but NOT at the beginning (or at the end) of a word (the "r" in the beginning is making sure that the string is being treated as a "raw string")	<code>r"\Bain"</code> <code>r"ain\B"</code>
<code>\d</code>	Returns a match where the string contains digits (numbers from 0-9)	<code>"\d"</code>
<code>\D</code>	Returns a match where the string DOES NOT contain digits	<code>"\D"</code>

Special Sequences

- A special sequence is a `\` followed by one of the characters in the list below, and has a special meaning:

<code>\D</code>	Returns a match where the string DOES NOT contain digits	<code>"\D"</code>
<code>\s</code>	Returns a match where the string contains a white space character	<code>"\s"</code>
<code>\S</code>	Returns a match where the string DOES NOT contain a white space character	<code>"\S"</code>
<code>\w</code>	Returns a match where the string contains any word characters (characters from a to Z, digits from 0-9, and the underscore <code>_</code> character)	<code>"\w"</code>
<code>\W</code>	Returns a match where the string DOES NOT contain any word characters	<code>"\W"</code>
<code>\Z</code>	Returns a match if the specified characters are at the end of the string	<code>"Spain\Z"</code>

Special Sequences

- **\A**

```
>> import re
>> txt = "The rain in Spain"
>> #Check if the string starts with "The":
>> x = re.findall("\AThe", txt)
>> print(x)
>> if x:
>>     print("Yes, there is a match!")
>> else:
>>     print("No match")
```

Output: ['The']
Yes, there is a match!

Special Sequences

- **\b**

```
>> import re
>> txt = "The rain in Spain"
>> #Check if "ain" is present at the end of a WORD:
>> x = re.findall(r"ain\b", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Special Sequences

- **\b**

```
>> import re
>> txt = "The rain in Spain"
>> #Check if "ain" is present at the beginning of a WORD:
>> x = re.findall(r"\bain", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Special Sequences

- **\B**

```
>> import re
>> txt = "The rain in Spain"
>> #Check if "ain" is present, but NOT at the beginning of a word:
>> x = re.findall(r"\Bain", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Special Sequences

- `\d`

```
>> import re
>> txt = "The rain in Spain"
>> #Check if the string contains any digits (numbers from 0-9):
>> x = re.findall("\d", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Special Sequences

- `\D`

```
>> import re
>> txt = "The rain in Spain"
>> #Return a match at every no-digit character:
>> x = re.findall("\D", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Special Sequences

- `\s`

```
>> import re
>> txt = "The rain in Spain"
>> #Return a match at every white-space character:
>> x = re.findall("\s", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Special Sequences

- `\S`

```
>> import re
>> txt = "The rain in Spain"
>> #Return a match at every NON white-space character:
>> x = re.findall("\S", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```


Special Sequences

- **\w**

```
>> import re
>> txt = "The rain in Spain"
>> #Return a match at every word character (characters from a to Z, digits from 0-9, and the
underscore _ character):
>> x = re.findall("\w", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Special Sequences

- **\W**

```
>> import re
>> txt = "The rain in Spain"
>> #Return a match at every NON word character (characters NOT between a and Z. Like "!",
"?" white-space etc.):
>> x = re.findall("\W", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Special Sequences

- `\Z`

```
>> import re
>> txt = "The rain in Spain"
>> #Check if the string ends with "Spain":
>> x = re.findall("Spain\Z", txt)
>> print(x)
>> if x:
>>     print("Yes, there is a match!")
>> else:
>>     print("No match")
```

Set

- A set is a set of characters inside a pair of square brackets `[]` with a special meaning:

Set	Description
<code>[arn]</code>	Returns a match where one of the specified characters (<code>a</code> , <code>r</code> , or <code>n</code>) are present
<code>[a-n]</code>	Returns a match for any lower case character, alphabetically between <code>a</code> and <code>n</code>
<code>[^arn]</code>	Returns a match for any character EXCEPT <code>a</code> , <code>r</code> , and <code>n</code>
<code>[0123]</code>	Returns a match where any of the specified digits (<code>0</code> , <code>1</code> , <code>2</code> , or <code>3</code>) are present
<code>[0-9]</code>	Returns a match for any digit between <code>0</code> and <code>9</code>
<code>[0-5][0-9]</code>	Returns a match for any two-digit numbers from <code>00</code> and <code>59</code>
<code>[a-zA-Z]</code>	Returns a match for any character alphabetically between <code>a</code> and <code>z</code> , lower case OR upper case
<code>[+]</code>	In sets, <code>+</code> , <code>*</code> , <code>.</code> , <code> </code> , <code>()</code> , <code>\$</code> , <code>{}</code> has no special meaning, so <code>[+]</code> means: return a match for any <code>+</code> character in the string

Set

- **[arn]**

```
>> import re
>> txt = "The rain in Spain"
>> #Check if the string has any a, r, or n characters:
>> x = re.findall("[arn]", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Output:
 ['r', 'a', 'n', 'n', 'a', 'n']
 Yes, there is at least one match!

Set

- **[a-n]**

```
>> import re
>> txt = "The rain in Spain"
>> #Check if the string has any characters between a and n:
>> x = re.findall("[a-n]", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Output:
 ['h', 'e', 'a', 'i', 'n', 'i', 'n', 'a', 'i', 'n']
 Yes, there is at least one match!

Set

- **[^arn]**

```
>> import re
>> txt = "The rain in Spain"
>> #Check if the string has other characters than a, r, or n:
>> x = re.findall("[^arn]", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Output:
 ['T', 'h', 'e', ' ', 'i', ' ', 'i', ' ', 'S', 'p', 'i']
 Yes, there is at least one match!

Set

- **[123]**

```
>> import re
>> txt = "The rain in Spain"
>> #Check if the string has any 0, 1, 2, or 3 digits:
>> x = re.findall("[0123]", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Output:
 []
 No match

Set

- [0-9]

```
>> import re
>> txt = "8 times before 11:45 AM"
>> #Check if the string has any digits:
>> x = re.findall("[0-9]", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Output:

['8', '1', '1', '4', '5']

Yes, there is at least one match!

Set

- [0=5][0-9]

```
>> import re
>> txt = "8 times before 11:45 AM"
>> #Check if the string has any characters from a to z lower case, and A to Z upper case:
>> x = re.findall("[a-zA-Z]", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Output:

['t', 'i', 'm', 'e', 's', 'b', 'e', 'f', 'o', 'r', 'e', 'A', 'M']

Yes, there is at least one match!

Set

- `[a-z][A-Z]`

```
>> import re
>> txt = "8 times before 11:45 AM"
>> #Check if the string has any characters from a to z lower case, and A to Z upper
case:
>> x = re.findall("[a-zA-Z]", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Output:

`['t', 'i', 'm', 'e', 's', 'b', 'e', 'f', 'o', 'r', 'e', 'A', 'M']`

Yes, there is at least one match!

Set

- `+`

```
>> import re
>> txt = "8 times before 11:45 AM"
>> #Check if the string has any + characters:
>> x = re.findall("[+]", txt)
>> print(x)
>> if x:
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Output:

`[]`

No match

Python RegEx

- Python has a module named **re** to work with regular expressions. To use it, we need to import the module.
- **import re**
- The module defines several functions and constants to work with RegEx.

Python RegEx

- **re.findall()**
- The **re.findall()** method returns a list of strings containing all matches.

Python RegEx

- **Example 1: re.findall()**
- # Program to extract numbers from a string

```
>> import re
>> string = 'hello 12 hi 89. Howdy 34'
>> pattern = '\d+'
>> result = re.findall(pattern, string)
>> print(result)
```

- # Output: ['12', '89', '34']
- If the pattern is not found, re.findall() returns an empty list.

Python RegEx

- **Example 2: re.findall()**

```
>> import re

>> #Return a list containing every occurrence of "ai":

>> txt = "The rain in Spain"
>> x = re.findall("ai", txt)
>> print(x)
```

Output:
['ai', 'ai']

Python RegEx

- **Example 3: re.findall()**

```
>> import re
>> txt = "The rain in Spain"
>> #Check if "Portugal" is in the string:
>> x = re.findall("Portugal", txt)
>> print(x)
>> if (x):
>>     print("Yes, there is at least one match!")
>> else:
>>     print("No match")
```

Output:
[]
No match

Python RegEx

- **re.split()**
- The **re.split** method splits the string where there is a match and returns a list of strings where the splits have occurred.

Python RegEx

- **Example 1: re.split()**

```
>> import re
>> string = 'Twelve:12 Eighty nine:89.'
>> pattern = '\d+'
>> result = re.split(pattern, string)
>> print(result)
```

- # Output: ['Twelve:', ' Eighty nine:', '.']
- If the pattern is not found, re.split() returns a list containing the original string.

Python RegEx

- **Example 2: re.split()**

```
>> import re
>> #Split the string at every white-space character:
>> txt = "The rain in Spain"
>> x = re.split("\s", txt)
>> print(x)
```

Output:

```
['The', 'rain', 'in', 'Spain']
```

Python RegEx

- **Example 3: re.split()**

```
>> import re
>> #Split the string at the first white-space character:
>> txt = "The rain in Spain"
>> x = re.split("\s", txt, 1)
>> print(x)
```

Output:

```
['The', 'rain in Spain']
```

Python RegEx

- You can pass **maxsplit** argument to the **re.split()** method. It's the maximum number of splits that will occur.

```
>> import re
>> string = 'Twelve:12 Eighty nine:89 Nine:9.'
>> pattern = '\d+'
>> # maxsplit = 1
>> # split only at the first occurrence
>> result = re.split(pattern, string, 1)
>> print(result)
```

- # Output: ['Twelve:', ' Eighty nine:89 Nine:9.']
- By the way, the default value of maxsplit is 0; meaning all possible splits.

Python RegEx

- **re.sub()**
- The syntax of re.sub() is:
 - re.sub(pattern, replace, string)
 - The method returns a string where matched occurrences are replaced with the content of replace variable.

Python RegEx

- **Example 1: re.sub()**

```
>> # Program to remove all whitespaces
>> import re

>> # multiline string
>> string = 'abc 12\
>> de 23 \n f45 6'

>> # matches all whitespace characters
>> pattern = '\s+'
```

Python RegEx

```
>> # empty string
>> replace = ""

>> new_string = re.sub(pattern, replace, string)
>> print(new_string)
```

```
>> # Output: abc12de23f456
```

If the pattern is not found, `re.sub()` returns the original string.

Python RegEx

- **Example 2: `re.sub()`**

```
>> import re
>> #Replace all white-space characters with the digit "9":
>> txt = "The rain in Spain"
>> x = re.sub("\s", "9", txt)
>> print(x)
```

Output:
The9rain9in9Spain

Python RegEx

- **Example 3: re.sub()**

```
>> import re
>> #Replace the first two occurrences of a white-space character with the digit 9:
>> txt = "The rain in Spain"
>> x = re.sub("\s", "9", txt, 2)
>> print(x)
```

Output:

The9rain9in Spain

Python RegEx

You can pass `count` as a fourth parameter to the `re.sub()` method. If omitted, it results to 0. This will replace all occurrences.

```
>> import re

>> # multiline string
>> string = 'abc 12\nde 23 \n f45 6'
```

Python RegEx

```
>> # matches all whitespace characters
>> pattern = '\s+'
>> replace = ''

>> new_string = re.sub(r'\s+', replace, string, 1)
>> print(new_string)
```

```
# Output:
# abc12de 23
# f45 6
```

Python RegEx

- **re.subn()**
- The **re.subn()** is similar to **re.sub()** except it returns a tuple of 2 items containing the new string and the number of substitutions made.
- **Example 4: re.subn()**

```
>> # Program to remove all whitespaces
>> import re

>> # multiline string
>> sstring = 'abc 12\nde 23 \n f45 6'
```

Python RegEx

```
>> # matches all whitespace characters
>> pattern = '\s+'

>> # empty string
>> replace = ''

>> new_string = re.subn(pattern, replace, string)
>> print(new_string)

>> # Output: ('abc12de23f456', 4)
```

Python RegEx

- **re.search()**
- The **re.search()** method takes two arguments: **a pattern and a string**. The method looks for the first location where the RegEx pattern produces a match with the string.
- If the search is successful, **re.search()** returns a match object; if not, it returns None.
- `match = re.search(pattern, str)`

Python RegEx

- **Example 1: re.search()**

```
>> import re
>> string = "Python is fun"
>> # check if 'Python' is at the beginning
>> match = re.search('\APython', string)

>> if match:
>> print("pattern found inside the string")
>> else:
>> print("pattern not found")
```

- # Output: pattern found inside the string

Python RegEx

- **Example 2: re.search()**

- import re
- txt = "The rain in Spain"
- x = re.search("\s", txt)
- print("The first white-space character is located in position:", x.start())
- Output:
- The first white-space character is located in position: 3

Python RegEx

- **Example 3: re.search()**

```
>> import re
```

```
>> txt = "The rain in Spain"
```

```
>> x = re.search("Portugal", txt)
```

```
>> print(x)
```

- Output:
- None

Python RegEx

- **Example 3: re.search()**

```
>> import re
```

```
>> txt = "The rain in Spain"
```

```
>> x = re.search("Portugal", txt)
```

```
>> print(x)
```

- Output:
- None

Python RegEx

- **Match object**
- You can get methods and attributes of a match object using `dir()` function.
- Some of the commonly used methods and attributes of match objects are:
 - `match.group()`
- The `group()` method returns the part of the string where there is a match.

Python RegEx

- **Example 2: Match object**

```
>> import re
>> string = '39801 356, 2102 1111'
>> # Three digit number followed by space followed by two digit number
>> pattern = '(\d{3}) (\d{2})'
>> # match variable contains a Match object.
>> match = re.search(pattern, string)
```

Python RegEx

```
>> if match:  
>>   print(match.group())  
>> else:  
>>   print("pattern not found")
```

- # Output: 801 35
- Here, match variable contains a match object.

Python RegEx

- Our pattern `(\d{3}) (\d{2})` has two subgroups `(\d{3})` and `(\d{2})`. You can get the part of the string of these parenthesized subgroups. Here's how:

```
>>> match.group(1)  
'801'  
  
>>> match.group(2)  
'35'  
>>> match.group(1, 2)  
('801', '35')  
  
>>> match.groups()  
('801', '35')
```

Python RegEx

- **Using r prefix before RegEx**
- When r or R prefix is used before a regular expression, it means raw string. For example, '\n' is a new line whereas r'\n' means two characters: a backslash \ followed by n.
- Backslash \ is used to escape various characters including all metacharacters. However, using r prefix makes \ treat as a normal character.

Python RegEx

- **Example 7: Raw string using r prefix**

```
>> import re  
  
>> string = '\n and \r are escape sequences.'  
  
>> result = re.findall(r'[\n\r]', string)  
>> print(result)
```

- **# Output: ['\n', '\r']**

Example of Regex

Example - Regex

- Import pandas as pd
- Import numpy as np

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

- `df = pd.read_csv('titanic_train.csv')`
- `df['Title'] = df['Name'].str.extract('([A-Za-z]+\.)', expand=False)`
- `df['Age'].fillna(df.groupby('title')['Age'].transform('mean'), inplace = True)`

Example 1 - Regex

- Import pandas as pd
- Import numpy as np

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

- `df = pd.read_csv('titanic_train.csv')`
- `df['Title'] = df['Name'].str.extract('([A-Za-z]+\.)', expand=False)`
- `df['Age'].fillna(df.groupby('title')['Age'].transform('mean'), inplace = True)`

Example 2 - Regex

- **Remove the unnecessary characters from columns**
- import pandas as pd
- import numpy as np
- `dfW = pd.read_csv('D:\\Teaching Subject\\Data Science\\Fall 2021\\Lectures\\Structuring and Regex Example\\weather_data.csv')`

	day	temperature	windspeed	event
0	1/1/2017	32	6us	Rain
1	1/4/2017	-9999	9	Sunny
2	1/5/2017	28	-7777	Snow
3	1/6/2017	-9999	7	NaN
4	1/7/2017	32 #	-7777	Rain
5	1/8/2017	-9999	-7777	Sunny
6	1/9/2017	-9999	-7777	NaN
7	1/10/2017	34FA	8yyy	Cloudy
8	1/11/2017	40	12	Sunny

Example 2 - Regex

- `dfW['temperature'].replace(['^0-9-'],'',inplace=True,regex=True)`
- **output**

	day	temperature	windspeed	event
0	1/1/2017	32	6us	Rain
1	1/4/2017	-9999	9	Sunny
2	1/5/2017	28	-7777	Snow
3	1/6/2017	-9999	7	NaN
4	1/7/2017	32	-7777	Rain
5	1/8/2017	-9999	-7777	Sunny
6	1/9/2017	-9999	-7777	NaN
7	1/10/2017	34	8yyy	Cloudy
8	1/11/2017	40	12	Sunny

Example 3 - Regex

- `pakistan_intellectual_capital`
- Practice of structuring problem. See in Jupyter notebook or PDF